

corVis: An R Package for Visualising Associations and Conditional Associations

by Amit Chinwan and Catherine Hurley

Abstract Correlation matrix displays are important tools to explore multivariate datasets. These displays with other measures of association can summarize interesting patterns to an analyst and assist them in framing questions while performing exploratory data analysis. In this paper, we present new visualisation techniques to visualise association between all the variable pairs in a dataset in a single plot, which is something existing displays lack. Also, we propose new methods to visualise relationship among variable pairs using conditioning. We use different layouts like matrix or linear for our displays. We use seriation in our displays which helps in highlighting interesting patterns easily. The R package corVis provides an implementation.

Section 1: Introduction

Correlation matrix display is a popular tool to visually explore correlations among variables while performing Exploratory Data Analysis (EDA) on a multivariate dataset. Popularized by [Friendly \(2002\)](#) as corgram, these displays are produced by first calculating the correlation among the variables and then plotting these calculated values in a matrix display. With effective ordering techniques, these displays quickly highlight variables which are highly correlated and an analyst interested in building a predictive model could use these displays to remove correlated variables and avoid multicollinearity.

The correlation displays are generally used with one of the Pearson's, Spearman's or Kendall's correlation coefficient and are therefore limited to quantitative variables. An analyst can use one-hot encoding of the qualitative variables in order to use these displays but will need to deal with the high dimensions as a result of the encoding. In addition to the dimensionality problem, it is not easy to assess the overall correlation when using the one-hot encoding. The existing methods to quickly explore association among qualitative variables in a dataset include using proportions or counts with different graphical displays like boxplots or barplots. Using association measures for qualitative pairs similar to correlation for quantitative pairs will help in summarizing the relationship, which then can be displayed like the correlation displays.

Tukey and Tukey ([Tukey and Tukey, 1985](#)) introduced scagnostics which are measures for scatterplots. Along with scagnostics, they proposed a scagnostics scatterplot matrix which is a visual display to explore and compare these measures for all the variable pairs in a dataset. By comparing multiple measures at once, the unusual variable pairs could be identified and looked at in more detail. In a similar manner, a display comparing association measures will help in finding interesting variable pairs. Many association measures have been proposed to summarize different types of relationships. The most commonly used measure is Pearson's correlation coefficient which captures any linear trend present between the variables. Other popular measures include Kendall's or Spearman's rank correlation coefficient which are non-parametric measures and looks for monotonic relationship. Distance correlation ([Székely et al., 2007](#)) is an important measure useful in exploring non-linear relationships. The information theory measure maximal information coefficient (MIC) ([Reshef et al., 2011](#)) is capable of summarizing complex relationships. With effective displaying techniques, the multiple measures of association provide a comparison tool that assist an analyst to reveal structure present in the data.

Small multiples (or Trellis display) is a simple yet powerful approach to compare partitions of data and understand multidimensional datasets ([Tufté, 1986](#)). The display is produced by splitting the data into groups by a conditioning variable and then plotting the data for each group. Such displays allow analysts to quickly infer about the impact of the conditioning variable. A similar idea applied to displays of association measures (correlation plot) will help uncover underlying patterns in the data. One such pattern is Simpson's paradox which can be detected by comparing Pearson's correlation for data at overall level versus individual levels of the conditioning variable.

In this paper, we propose extensions of the correlation plot and new visualizations which look at variables of mixed type, multiple association measures and conditional associations. These displays are implemented in the R package corVis. The next section provides a review of existing packages which deal with correlation displays and a quick background on association measures and the packages used for calculating them. Then we describe our approach to calculate the association measures, followed by visualizations of associations and conditional associations. We conclude with a summary and future work.

Section 2: Background

In this section we provide a brief review of existing packages used for correlation displays and association measure calculation.

Section 2.1: Literature Review on Correlation Displays

According to [Hills \(1969\)](#), “the first and sometimes only impression gained by looking at a large correlation matrix is its largeness”. To overcome this, [Murdoch and Chow \(1996\)](#) proposed a display for large correlation matrices which uses a matrix layout of ellipses where the parameters of the ellipses are scaled to the correlation values. [Friendly \(2002\)](#) expanded on this idea by rendering correlation values as shaded squares, bars, ellipses, or circular ‘pac-man’ symbols. The variables in the matrix displays were optionally ordered using the angular ordering of the first two eigen vectors of the correlation matrix. The ordering places highly-correlated pairs of variables nearby, making it easier to quickly identify groups of variables with high mutual correlation.

Nowadays, there are many R packages devoted to correlation visualisation. Table 1 provides a summary, listing the displays offered, and whether these extend to factor variables or mixed numeric-factor pairs.

The R package [corrplot](#) ([Wei and Simko, 2021](#)) provides an implementation of the methods in [Friendly \(2002\)](#). The package [corrr](#) ([Kuhn et al., 2020](#)) organises correlations as tidy data, so leveraging the data manipulation and visualisation tools of the [tidyverse](#) ([Wickham et al., 2019](#)). In addition to various matrix displays, the package offers network displays where line-thickness encodes correlation magnitude, with a filtering option to discard low-correlation edges.

The package [corrgrapher](#) ([Morgen and Biecek, 2020](#)) uses a network plot for exploring correlations, where the nodes close to each other have high correlation magnitude, edge thickness encodes the absolute correlation value and edge color indicates the sign of correlation. The package also handles mixed type variables by using association measures obtained as transformations of p -values obtained from Pearson’s correlation test in the case of two numeric variables, Kruskal’s test for numerical and factor variables, and a chi-squared test for two categorical variables.

The package [linkspotter](#) ([Samba, 2020](#)) offers a variety of association measures (distance correlation, MIC, maximum normalized mutual information) in addition to correlation, where the measure used depends on whether the variables are both numerical, categorical or mixed. The results are visualized in a network plot, which may be packaged into an interactive shiny application.

Our own package [corVis](#) offers a variety of displays, and has new features not available elsewhere, in particular simultaneous display of multiple association measures, and association displays stratified by levels of a grouping variable. This will be described in the following sections.

There have been other extensions to correlation displays which are useful when dealing with high dimensional datasets. [Hills \(1969\)](#) proposed a QQ plot of the z -transform of the entries of the correlation matrix to discover correlation coefficients too large to come from a normal distribution with mean zero. [Buja et al. \(2016\)](#) proposed Association Navigator which is an interactive visualization tool for large correlation matrices with upto 2000 variables. The R package [scorrplot](#) ([Gerber, 2022](#)) produces an interactive scatterplot for exploring pairwise correlations in a large dataset by projecting variables as points and encoding the correlations as space between these points. The package provides a functionality to update variable of interest which creates tour of the correlation space between different projections of the data.

The R package [correlationfunnel](#) offers a novel display which assists in feature selection in a setting with a single response and many predictor variables. All numeric variables including the response are binned. All (now categorical) variables in the resulting dataset are one-hot encoded and Pearson’s correlation calculated with the response categories. The correlations are visualised in a dot-plot display, where predictors are ordered by maximum correlation magnitude. Correlations between one-hot encoded variables are challenging to interpret, especially as the number of levels increase. In [corVis](#) we offer a similar dot-plot display, but showing multiple correlation or association measures, or alternatively measures stratified by a grouping variable.

Section 2.2: Literature Review on Association Measures

An association measure is defined as a numerical summary quantifying the relationship between two or more variables. The measure is called symmetric if its value is invariant to the choice of independent or dependent variable during the calculation. For example, Pearson’s correlation coefficient summarizes the strength and direction of the linear relationship present between two numeric variables in the range $[-1, 1]$ and is symmetric. Kendall’s or Spearman’s rank correlation coefficient are other popular

Table 1: List of the R packages dealing with correlation or correlation displays with information on whether the plots display multiple measures, conditional display of measures and mixed variables in a single plot

Package	Display	MixedVariables
corrplot	heatmap	
corr	heatmap/network	
corrgrapher	network	
linkspotter	network	Yes
correlation	heatmap/network	
corVis	heatmap/matrix/linear	Yes




Figure 1: Multiple association measures for simulated linear and non-linear pattern. The first row of the plot shows a linear pattern, the value of Pearson’s correlation, distance correlation and MIC. The second row of the plot shows a non-linear pattern along with the values for association measures. All three association measures show a high value for the linear relationship and hence are useful for linear patterns. For non-linear pattern, distance correlation and MIC are more suitable measures than Pearson’s correlation.

measures which assess monotonic relationship among two numeric variables in interval $[-1, 1]$ and are symmetric measures.

Pearson’s correlation is the most popular association measure because of its easier calculation and interpretability but its limitations such as influence of outliers on its value and measuring only linear dependencies makes it a less useful measure of association overall. The recent measures such as distance correlation (Székely et al., 2007) and MIC (Reshef et al., 2011) overcome these limitations and are more suitable for datasets with both linear and non-linear patterns.

Figure 1 shows a plot of simulated linear and non-linear patterns. The first row shows a linear relationship along with the value for measures such as Pearson’s correlation, distance correlation and maximal information coefficient respectively for the pattern. In a similar manner, the second row shows a quadratic relationship and its values for Pearson’s correlation, distance correlation and maximal information coefficient respectively. It is clearly evident from Figure 1 that all the three measures summarizes the pattern with linear relationship quite well. For the non-linear pattern, distance correlation and MIC are better in detecting underlying relationship than Pearson’s correlation. This suggests that association measures such as distance correlation, MIC and other measures should be used along with Pearson’s correlation for exploring relationships among variables in the datasets where there is no prior knowledge about the possible patterns in the data.

The distance correlation coefficient (Székely et al., 2007) is an association measure which looks for any relationship among two numeric variables using the distances between observations of these variables and summarizes the relationship in $[0, 1]$. The distance correlation is 0 only when the variables are independent and is a symmetric measure.

The maximal information coefficient (MIC) (Reshef et al., 2011) is an information theory measure which uses mutual information among the two variables for its calculation. The main idea is to find a grid out of possible grids on a scatterplot of two numeric variables, in order to discretize the variables, which maximises the mutual information for the two variables. A normalisation technique is used to make the mutual information from different grids comparable. Referred as ‘a correlation of 21st century’ [speed2011correlation], MIC is capable of summarizing different types of relationships, not just linear or monotonous, between numeric variables and is in range $[0, 1]$. Reshef et al. (2011) used MIC and other related statistics to explore pairwise relationships in large data sets such as major-league baseball, gene expression, global health, and the human gut microbiota.

In addition to association measures for numeric variables, association measures for ordinal, nominal and mixed variable pairs are useful in exploring a multivariate dataset. We now give an overview of available association measures for other variable types.

Agresti (2010) provides an overview of the association measures which are used for exploring association between ordinal variables. Kendall’s tau-b (Kendall, 1945) is an association measure useful in summarizing the relationship between two ordinal variables in the range $[-1, 1]$. It is a relatively stable measure than Goodman and Kruskal’s gamma with respect to the changes in categories of any variable i.e. if two categories are merged to make a single category. The polychoric correlation (Olsson, 1979) measures the correlation between two ordinal variables by assuming two normally distributed latent variables for a contingency table of two ordinal variables and summarizes the association in $[-1, 1]$.

The association measures for the case of nominal pair of variables should be invariant to the order in which the categories appear. Pearson’s contingency coefficient uses the χ^2 value from the Pearson’s χ^2 test for independence and is a useful measure to summarize the association between two nominal variables in $[0, 1]$. Another measure for nominal variable pair is the Uncertainty coefficient (Theil, 1970) measuring the proportion of uncertainty in one variable which is explained by the other variable.

Section 3: Introducing corVis

corVis is an R package which calculates measures of association for every variable pair in a dataset and provides visualizations for displaying associations. Most of the existing correlation displays are limited to numeric pairs of variables. This package extends these displays to every variable pair. The main goal of our work is to propose displays for multiple association measures and conditional associations display which are useful for uncovering interesting patterns in the data. This will help in identifying variable pairs which shows a type of relationship or pattern in a dataset with large number of variables.

Consider Figure 2 for the motivation behind conditional association displays. The plot on the left shows a positive linear association between y and x and has a positive Pearson’s correlation of value 0.603. The plot on the right shows the disaggregated data by the group variable and it is clearly evident that for each group there is a negative linear relationship between y and x . This is an example of Simpson’s paradox and is one of the patterns which are discoverable when conditional association displays are used.

While designing these displays we considered matrix and linear layouts. A matrix layout simplifies the effort in finding variables, and different measures may be displayed on the upper and lower diagonal. Linear layouts are more space-efficient than matrix plots, but looking for variables or variable pairs is more challenging.

Table 2 provides a list of the functions available in the package which are useful for calculating association measures among variable pairs and visualising these associations using novel displays. Section 4 provides a detailed description on functions used for calculating association measures in the package. Section 5 illustrates the use of visualising functions in table 2 for displaying pairwise association and conditional association.

Section 4: corVis: Calculating Association

This section describes the calculation of association measures in our package **corVis**. The package provides a standard interface for calculating a collection of various measures of association which quantifies the relationship between two variables. The association measures available in the package




Figure 2: An example plot showing importance of conditional displays

Table 2: List of functions in corVis package

Function	Usage	Description
<code>calc_assoc</code>	Calculation	Calculates association measures
<code>association_heatmap</code>	Visualization	Conventional correlation matrix plot
<code>pairwise_2d_plot</code>	Visualization	Conditional and multiple measures matrix plot
<code>pairwise_1d_plot</code>	Visualization	Conditional and multiple measures linear plot

Table 3: List of the functions available in the package for calculating different association measures along with the packages used for calculation.

Function	X	Y	from	symmetric
tbl_cor	numerical	numerical	stats::cor	Y
tbl_dcor	numerical	numerical	energy::dcor2d	Y
tbl_mine	numerical	numerical	minerva::mine	Y
tbl_polycor	ordinal	ordinal	polycor::polychor	Y
tbl_tau	ordinal	ordinal	DescTools::KendalTauA,B,C,W	Y
tbl_gkGamma	ordinal	ordinal	DescTools::GoodmanKruskalGamma	Y
tbl_gkTau	nominal	nominal	DescTools::GoodmanKruskalTau	N
tbl_uncertainty	nominal	nominal	DescTools::UncertCoef	Y
tbl_chi	nominal	nominal	DescTools::ContCoef	Y
tbl_cancor	nominal/numerical	nominal/numerical	corVis	Y
tbl_nmi	nominal	nominal	corVis	Y
tbl_easy	nominal/numerical	nominal/numerical	correlation::correlation	Y

are not limited to numeric variables and are used with nominal, ordinal and mixed variable pairs as well. The package also provides a functionality for handling missing value or NA while calculating the association measures.

Table 3 lists different functions provided in the package to calculate various measures of association. The Function column represents the function name used to calculate measure(s) of associations in this package. The typeX and typeY columns provide the information on types of variables which can be used with the corresponding functions. The X or Y variable is one of the numeric, nominal, ordinal or any type. The from column corresponds to the package functions used to calculate the association measures by the function under Function. The symmetric column represents if the measure is symmetric i.e. if the value of measure is same regardless of the order of variables. The last column provides the range of values for these measures. The function tbl_easy calculates association measures available in the R package [correlation](#) and is suitable for different variable types. The functions in Table 3 with corVis entries under from column calculate the association measures which have been implemented in this package.

For numeric pairs of variables, this package provides a range of association measures. The popular correlation coefficients like Pearson's or Spearman's or Kendall's are calculated using tbl_cor function. The measures such as distance correlation or MIC are calculated using tbl_dcor or tbl_mine respectively. The association measures available in the package for the ordinal pairs of variables are polychoric correlation and Kendall's coefficients which are calculated using tbl_polycor or tbl_tau respectively. For nominal pairs of variables, the functions like tbl_gkTau, tbl_gkGamma, tbl_uncertainty, tbl_chi, tbl_cancor are used for exploring association among the variables.

The function tbl_cancor calculates a measure of association based on canonical correlations for mixed pairs of variables. Nominal variables are converted into sets of dummy variables, which are then assigned scored to find the maximal correlation. For two numeric variables this measure is identical to absolute correlation, for two factors the correlation is identical to that obtained from correspondence analysis.

The functions listed in 3 for calculating association measures provide a functionality for handling missing value or NA in the dataset. Each of these functions either have a handle.na argument or have package functions which automatically uses pairwise complete observations for taking care of missing values present in the data.

Calculating association for a single type of variable pairs

We have a function which creates a tibble structure for the variable pairs in a dataset along with calculated association measure. The package contains various functions (shown in Table 3) for different association measures in the form tbl_* to calculate them. For example, in order to calculate distance correlation for numeric pair of variables in a dataset, the function tbl_dcor is used.

```
df <- iris
```

```
distance <- tbl_dcor(df)
```



```
distance

#> # A tibble: 6 x 4
#>   x          y          measure measure_type
#>   <chr>      <chr>      <dbl> <chr>
#> 1 Sepal.Width Sepal.Length 0.311 dcor
#> 2 Petal.Length Sepal.Length 0.859 dcor
#> 3 Petal.Width  Sepal.Length 0.827 dcor
#> 4 Petal.Length Sepal.Width 0.542 dcor
#> 5 Petal.Width  Sepal.Width 0.513 dcor
#> 6 Petal.Width  Petal.Length 0.974 dcor
```

The tibble output for the functions mentioned in Table 3 has the following structure:

- x and y representing a pair of variables
- measure representing the calculated value for association measure
- measure_type representing the association measure calculated for x and y pair.

Calculating association measures for whole dataset

calc_assoc is used to calculate association measures for all the variable pairs in the dataset at once in a tibble structure. The variable pairs in the output are unique pairs and a subset of all the pairs of variables in a dataset where $x \neq y$. Because of the tidy structure of the output, the data manipulation and visualisation tools of [tidyverse](#) (Wickham et al., 2019) are applicable to and are useful for further exploration of pairwise associations. In addition to tibble structure, the output also has pairwise and data.frame class which are important class attributes for producing visual summaries in this package.

The function calc_assoc has a types argument which is basically a tibble of the association measure to be calculated for different variable pairs. The default tibble of measures is default_assoc() which calculates Pearson's correlation if both the variables are numeric, Kendall's tau-b if both the variables are ordinal, canonical correlation if one is factor and other is numeric and canonical correlation for the rest of the variable pairs.

```
default_measures <- default_assoc()
default_measures

#> # A tibble: 4 x 4
#>   funName    typeX    typeY    argList
#>   <chr>      <chr>    <chr>    <list>
#> 1 tbl_cor    numeric numeric <NULL>
#> 2 tbl_tau    ordered ordered <NULL>
#> 3 tbl_cancor factor  numeric <NULL>
#> 4 tbl_cancor other   other   <NULL>

iris_assoc <- calc_assoc(df, types = default_assoc())
iris_assoc

#> # A tibble: 10 x 4
#>   x          y          measure measure_type
#>   <chr>      <chr>      <dbl> <chr>
#> 1 Sepal.Width Sepal.Length -0.118 pearson
#> 2 Petal.Length Sepal.Length 0.872 pearson
#> 3 Petal.Width  Sepal.Length 0.818 pearson
#> 4 Species      Sepal.Length 0.787 cancel
#> 5 Petal.Length Sepal.Width -0.428 pearson
#> 6 Petal.Width  Sepal.Width -0.366 pearson
#> 7 Species      Sepal.Width 0.633 cancel
#> 8 Petal.Width  Petal.Length 0.963 pearson
#> 9 Species      Petal.Length 0.970 cancel
#> 10 Species     Petal.Width 0.964 cancel

class(iris_assoc)

#> [1] "pairwise" "tbl_df" "tbl" "data.frame"
```

The default tibble of measures is updated using the `update_assoc` function which has arguments for updating the `tbl_*` functions to calculate association measures depending on the type variable pair in the dataset and a method for `tbl_*` functions which calculates more than one measure. The `update_assoc` function has an argument `default` which has the `default_assoc()` tibble as its default value and is useful when `tbl_*` functions need to be updated for a few types of variable pairs.

```
updated_assoc <- update_assoc(default=default_assoc(),
                             num_pair = "tbl_cor",
                             num_pair_argList = "spearman",
                             mixed_pair = "tbl_cancor",
                             other_pair = "tbl_nmi")

updated_assoc

#> # A tibble: 4 x 4
#>   funName   typeX   typeY   argList
#>   <chr>    <chr>  <chr>  <list>
#> 1 tbl_cor   numeric numeric <chr [1]>
#> 2 tbl_tau   ordered ordered <NULL>
#> 3 tbl_cancor factor  numeric <NULL>
#> 4 tbl_nmi   other   other   <NULL>
```

`calc_assoc` also has a `handle.na` argument for handling the NA or missing values which is fed into the `tbl_*` functions used with the `types` argument for different types of variable pairs. The default value is set to TRUE for using pairwise complete observations for calculating a measure of association between two variables.

If a user is interested in calculating multiple association measures for a type of variable pair, it can be done by using the `calc_assoc` and `update_assoc` together for calculating different association measures and then merging the output tibbles.

```
updated_iris_assoc <- calc_assoc(df, types = updated_assoc)
updated_iris_assoc

#> # A tibble: 10 x 4
#>       x             y      measure measure_type
#>   <chr>        <chr>    <dbl> <chr>
#> 1 Sepal.Width Sepal.Length -0.167 spearman
#> 2 Petal.Length Sepal.Length  0.882 spearman
#> 3 Petal.Width  Sepal.Length  0.834 spearman
#> 4 Species      Sepal.Length  0.787 cancel
#> 5 Petal.Length Sepal.Width  -0.310 spearman
#> 6 Petal.Width  Sepal.Width  -0.289 spearman
#> 7 Species      Sepal.Width   0.633 cancel
#> 8 Petal.Width  Petal.Length  0.938 spearman
#> 9 Species      Petal.Length  0.970 cancel
#> 10 Species     Petal.Width   0.964 cancel
```

Calculating conditional association

`calc_assoc` is also used to calculate association measures for all the variable pairs at different levels of a categorical variable. This helps in exploring the conditional associations and find out the differences between the groups of the conditioning variable. The function has a `by` argument which is used as the grouping variable and needs to be categorical.

```
iris_assoc_by <- calc_assoc_by(df, by = "Species")
iris_assoc_by

#> # A tibble: 24 x 5
#>       x             y      measure measure_type by
#>   <chr>        <chr>    <dbl> <chr>    <fct>
#> 1 Sepal.Width Sepal.Length  0.743 pearson  setosa
#> 2 Petal.Length Sepal.Length  0.267 pearson  setosa
#> 3 Petal.Width  Sepal.Length  0.278 pearson  setosa
#> 4 Petal.Length Sepal.Width   0.178 pearson  setosa
#> 5 Petal.Width  Sepal.Width   0.233 pearson  setosa
#> 6 Petal.Width  Petal.Length  0.332 pearson  setosa
```



```
#> 7 Sepal.Width Sepal.Length 0.526 pearson versicolor
#> 8 Petal.Length Sepal.Length 0.754 pearson versicolor
#> 9 Petal.Width Sepal.Length 0.546 pearson versicolor
#> 10 Petal.Length Sepal.Width 0.561 pearson versicolor
#> # ... with 14 more rows
```

By default, the function `calc_assoc` calculates the association measures for all the variable pairs at different levels of the grouping variable and the pairwise association measures for the ungrouped data (overall) when used with the `by` argument. This behavior can be changed by setting `include.overall = FALSE` argument to `FALSE`.

```
iris_assoc_by <- calc_assoc_by(df, by = "Species", include.overall = FALSE)
iris_assoc_by
```

```
#> # A tibble: 18 x 5
#>   x          y      measure measure_type by
#>   <chr>      <chr>      <dbl> <chr>      <fct>
#> 1 Sepal.Width Sepal.Length 0.743 pearson setosa
#> 2 Petal.Length Sepal.Length 0.267 pearson setosa
#> 3 Petal.Width Sepal.Length 0.278 pearson setosa
#> 4 Petal.Length Sepal.Width 0.178 pearson setosa
#> 5 Petal.Width Sepal.Width 0.233 pearson setosa
#> 6 Petal.Width Petal.Length 0.332 pearson setosa
#> 7 Sepal.Width Sepal.Length 0.526 pearson versicolor
#> 8 Petal.Length Sepal.Length 0.754 pearson versicolor
#> 9 Petal.Width Sepal.Length 0.546 pearson versicolor
#> 10 Petal.Length Sepal.Width 0.561 pearson versicolor
#> 11 Petal.Width Sepal.Width 0.664 pearson versicolor
#> 12 Petal.Width Petal.Length 0.787 pearson versicolor
#> 13 Sepal.Width Sepal.Length 0.457 pearson virginica
#> 14 Petal.Length Sepal.Length 0.864 pearson virginica
#> 15 Petal.Width Sepal.Length 0.281 pearson virginica
#> 16 Petal.Length Sepal.Width 0.401 pearson virginica
#> 17 Petal.Width Sepal.Width 0.538 pearson virginica
#> 18 Petal.Width Petal.Length 0.322 pearson virginica
```

The tibble output in the conditional setting has a similar structure as `calc_assoc` used with no `by` argument. When used with the `by` argument, an additional `by` column representing the levels of the categorical variable is added in the tibble output. The `x` and `y` variables in the output are repeated for every level of `by` variable. In order to have multiple `by` variables, the function `calc_assoc` is used multiple times with a different `by` variable each time and then the multiple outputs are binded row wise. For calculating multiple measures for a specific variable type, one can use `update_assoc` with `calc_assoc` and then can merge these multiple tibble outputs.

Section 5: corVis: Visualising Association

This section provides a detailed description of the novel visualisation techniques proposed in the package `corVis`. These methods display association and conditional association for every variable pair in a dataset in a single plot and show multiple bivariate measures of association simultaneously.

We use Steel Industry Energy consumption dataset providing information on energy consumption by a steel company in South Korea from the (Dua and Graff, 2017) to provide illustrative examples. Table 4 provides a brief description of set of variables from this dataset used for analysis.

Association Matrix plot

The function `association_heatmap` is used to display a matrix layout with association for variable pairs in the dataset. The display is similar to existing correlation matrix plots but with every variable pair in the dataset. This function `association_heatmap` takes the calculated measures of association by `calc_assoc` function as input and outputs a matrix display by rendering the magnitude of association measures with a color. The function has `lassoc` and `uassoc` arguments for a tibble of association measures for the lower triangle and the upper triangle of the matrix display respectively. The `uassoc` argument is `NULL` by default and uses the same tibble input as used by `lassoc` if not changed. The argument `var_order` is used for ordering or seriating the matrix display such that highly-associated

Table 4: Variable description of the Steel Industry Energy consumption dataset along with the types of variables

Variable	Description	VariableType
Usage_kWh	Energy Consumption in kWh	numerical
LgCRP	Lagging Current reactive power in kVarh	numerical
LdCRP	Leading Current reactive power in kVarh	numerical
CO2.tCO2.	CO2 emission in ppm	numerical
LgCPF	Lagging Current power factor	numerical
LdCPF	Leading Current power factor	numerical
NSM	Number of Seconds from midnight	numerical
WeekStatus	Week status (0 for weekend, 1 for weekday)	nominal
Day_of_week	Day of the week	nominal
Load_Type	Load type (Light Load, Medium Load or Maximum Load)	ordinal

variables are placed nearby and are easier to identify. We use average linkage hierarchical clustering method as the default method for ordering the variables. The function also has a `limits` argument specifying the limit of the color scale.

Figure 3 shows an example of the association matrix display for every variable pair in the energy consumption for steel industry dataset. The cells along the diagonal of the matrix display show the variables present in the dataset. Every off diagonal cell is colored using a divergent color scale with limits $[-1, 1]$ representing the value of association measure between two variables. The plot shows Pearson's correlation for the numeric pair of variables and canonical correlation for mixed and nominal variable pairs.

```
df <- steelindustry[sample(1:nrow(steelindustry),5000),] # taking a sample as original dataset has 35k entries
#df$al <- as.ordered(df$al)
#df$su <- as.ordered(df$su)
```

```
assoc_steel <- calc_assoc(df)
association_heatmap(assoc_steel)
```

For numeric pairs, Figure 3 shows a high positive Pearson's correlation among Usage_kWh and Co2.tCO2. and a high negative Pearson's correlation between LdCRP and LdCPF. Figure 4 shows scatterplot for these two pairs of variables and is evident that these pairs seem to be linearly associated.

Figure 3 shows a high canonical correlation between Load_Type and Usage_kWh, and, Load_Type and NSM which are mixed variable pairs in the dataset. These variable pairs are further looked into using boxplots. Figure 5 shows boxplot for these variable pairs and is evident there is an association among these variable pairs.

For nominal variable pairs, Figure 3 shows a high canonical correlation among WeekStatus and Day_of_week only. Because both variables are nested variables, there is a high association among them and is picked by canonical correlation.

Multiple Association Measures Plot

The function `pairwise_1d_compare` is used to display a linear layout for comparing multiple association measures calculated for variable pairs in the dataset. `pairwise_1d_compare` takes the calculated association measures as input and outputs a heatmap display by plotting all the variable pairs in the dataset on the Y-axis and the association measures on the X-axis and encoding the magnitude of association measures with a color.

The `assoc` argument of the function `pairwise_1d_compare` represents a tibble of calculated multiple association measures. These measures are calculated using `tbl_*` or `calc_assoc` functions and are binded row wise to create a tibble structure with multiple measures. The function also has a `var_order` argument which is set to `max_diff` by default and orders the variable pairs on the basis of maximum difference among the magnitude of measures.

Figure 6 shows the output of `pairwise_1d_compare` comparing association measures such as canonical correlation, MIC, Pearson's correlation, Spearman's correlation and distance correlation for all the variable pairs in the dataset. The cells where association measures are not defined for specific




Figure 3: Association matrix display for energy consumption data showing Pearson’s correlation for numeric variable pairs, canonical correlation for mixed variable pairs and categorical variable pairs.




Figure 4: Scatterplot for variables Usage in kWh and CO2 in ppm, and, LdCRP and LdCPF showing a high positive and negative Pearson’s correlation respectively.




Figure 5: Boxplot for variables Load type and Usage in kWh, and, Load type and NSM showing a high canonical correlation.

variable pairs have been colored grey. An interesting variable pair evident in the plot is LgCPF and LgCRP which has a low value for all the measures except MIC. On looking at the scatterplot (Figure 7) of LgCPF and LgCRP, it seems MIC is able to detect the non-linear pattern quite well in this case compared to the other association measures.

```
pearson_steel <- tbl_cor(df, method = "pearson")
spearman_steel <- tbl_cor(df, method = "spearman")
distance_steel <- tbl_dcor(df)
mic_steel <- tbl_mine(df)
cancor_steel <- tbl_cancor(df)
```

```
compare_assoc_steel <- rbind(pearson_steel, spearman_steel, distance_steel, mic_steel, cancor_steel)
compare_assoc_steel$measure_type <- factor(compare_assoc_steel$measure_type, levels = c("cancor", "mic", "dcor", "pearson", "spearman"))
pairwise_1d_compare(compare_assoc_steel)
```

Conditional Association Plot

The function `pairwise_2d_plot` is used to display a matrix layout of the conditional association for variable pairs in the dataset. The display is produced by splitting the data by a partitioning variable and calculating association for the variable pairs at each level of partitioning variable using `calc_assoc` function with conditioning variable as the `by` argument. The calculated association measures are then displayed using bars in a matrix plot. The height and color of the bars are coded with the value of association measure and the level of the partitioning variable respectively. These displays are efficient for discovering variable pair with high differences among the levels of partitioning variable in the data.

The measures of association calculated for every variable pair at every level of conditioning serve as input to the `pairwise_2d_plot` function. The `lassoc` and `uassoc` arguments look for a tibble of association measures for the lower triangle and the upper triangle of the matrix display respectively. The `uassoc` argument is set to `NULL` by default and uses the same tibble input as used by `lassoc` if not updated. The argument `group_var` is responsible for the grouping variable when plotting the bars.




Figure 6: Comparing multiple association measures using a linear layout. The display has variable pairs on the Y-axis and association measures on the X-axis. The cell corresponding to a variable pair and an association measure has been colored grey if the measure is not defined for corresponding pair.




Figure 7: Scatterplot for numeric variable pair LgCPF and LgCRP with high difference among the magnitude of association measures

The default value by uses the conditioning variable and measure_type is useful for displaying bars with height and color of the bars coded by the value of association measure and the type of association measure respectively.

Figure 8 shows a conditional association plot for the Steel industry energy consumption data. Each cell corresponding to a variable pair shows seven bars which correspond to the association measure (MIC for numeric pair and canonical correlation coefficient for mixed and nominal variable pairs) calculated at the levels of conditioning variable Day_of_week. The dotted line represents the overall association measure. The plot indicates that there is a high difference in the measure (MIC) value for the variable pair LgCRP and Usage_kWh for Saturday and Sunday, compared to other days and overall value.

```
updated_assoc <- update_assoc(default=default_assoc(),
                              num_pair = "tbl_mine",
                              mixed_pair = "tbl_cancor")
cond_assoc_steel <- calc_assoc_by(df, by="Day_of_week", types = updated_assoc)
pairwise_2d_plot(cond_assoc_steel)
```

Figure 9 shows the scatterplot for the variable pair LgCPR and Usage_kWh faceted by the conditioning variable Day_of_week. The scatterplot for Sunday shows a weaker association than other days of the week and could be investigated in more detail. It is possible to filter out more variable pairs from the conditional association plot having high difference for association measure among the groups and further investigate these pairs.

We also use linear layouts for displaying conditional association in the package. The function pairwise_1d_plot is used for displaying a linear layout of the conditional association for variable pairs in the dataset. The association measures are calculated for every variable pair at each level of partitioning variable using calc_assoc function with conditioning variable as the by argument.

The measures are then displayed using a dotplot where color of the dots are coded by the level of the partitioning variable and the variable pairs are ordered by absolute maximum value of association measure for each of the pair of variable. These displays are also efficient for discovering differences among the levels of partitioning variable in the data. With the linear layouts it is easier to omit less relevant pairs of variables by filtering the variables pairs having a higher value for association measures than a threshold.

The measures of association calculated for every variable pair at every level of conditioning serve as input to the pairwise_1d_plot function. The assoc argument uses a tibble of association measures

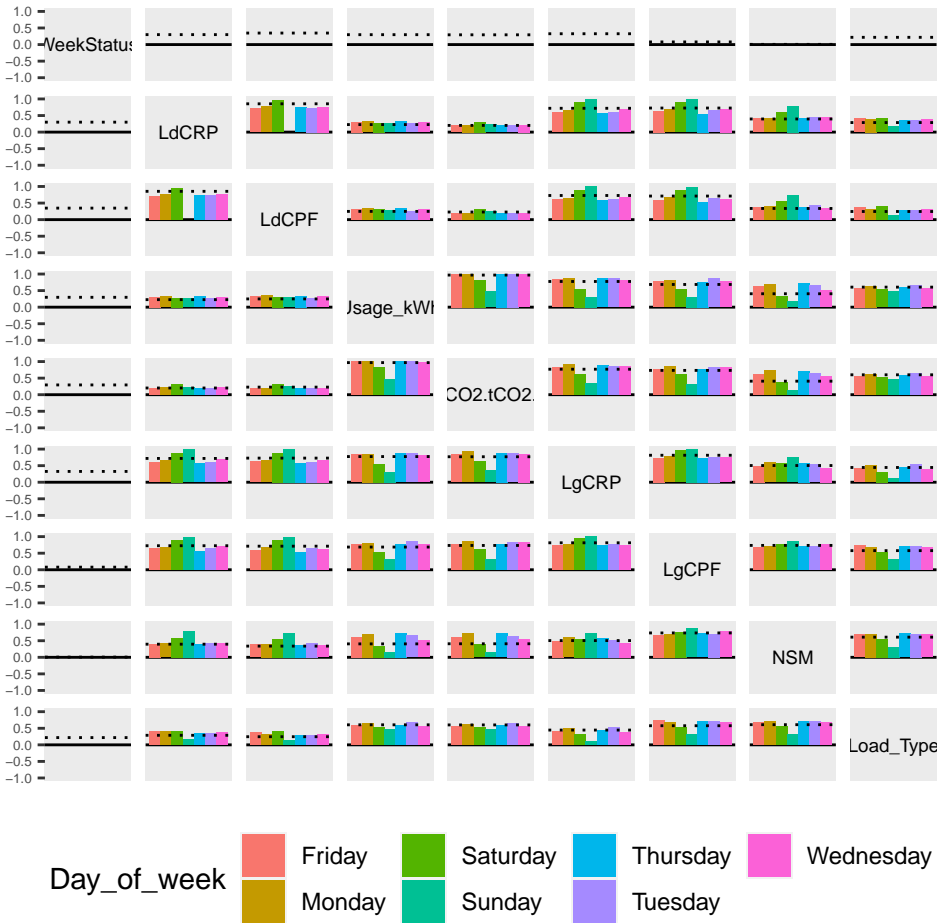


Figure 8: Conditional Association plot for energy consumption data showing MIC for numeric pairs and canonical correlation for nominal or mixed pairs. The bars in each cell represent the value for association measure colored by the conditioning variable Day of week. The dotted line in each cell represents overall value of the association measure.

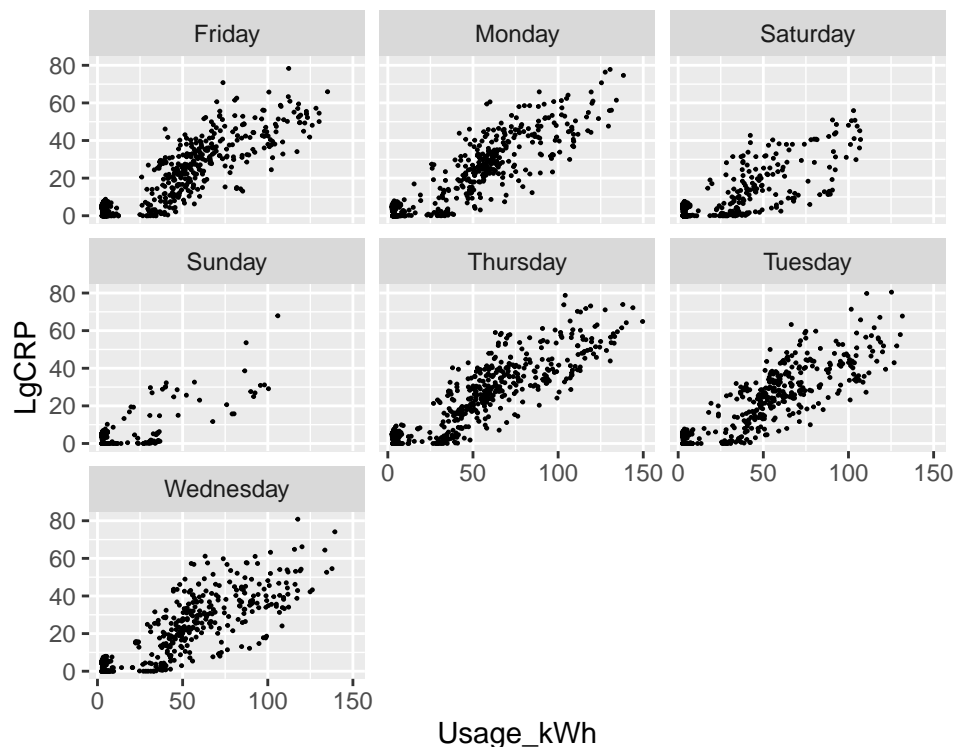


Figure 9: Scatterplot for variable pair LgCPR and Usage in kWh splitted by conditioning variable Day of week.

calculated using `calc_assoc` function with a `by` variable. The argument `group_var` is responsible for the grouping variable when plotting the dots. The default value `by` uses the conditioning variable and `measure_type` is useful for displaying dots with color of the dots coded by the type of association measure. The `var_order` argument is responsible for the ordering of variable pairs in the display. If set to `default` variable pairs are ordered alphabetically and are ordered by absolute maximum value of association measure for every variable pair when set to `max_diff`.

Figure 10 shows a funnel-like linear display for conditional association measures with all the variable pairs on the y-axis, the value of association measure on x-axis and color of the points representing the level of the grouping variable. The linear layout becomes more useful over the matrix layout when the number of variables and number of levels of grouping variable are high.

```
pairwise_1d_plot(cond_assoc_steel)
```

Section 5: Discussion

We use multiple association measures in a single display for different variable pairs which serves as a comparison tool while exploring association in a dataset and assist in identifying unusual variable pairs. These multiple measures can be displayed in a scatterplot matrix similar to what [Tukey and Tukey \(1985\)](#) proposed. They suggested that scatterplot matrix of the scagnostics measures, which are measures summarizing a scatterplot, can be used to identify unusual scatterplots or variable pairs. [Wilkinson et al. \(2005\)](#) used this idea with their graph-theoretic scagnostic measures to highlight unusual scatterplots. Similarly, [Kuhn et al. \(2013\)](#) have used this idea in a predictive modeling context. They have produced a scatterplot matrix of the measures between the response and continuous predictors such as Pearson's correlation coefficient, pseudo- R^2 from the locally weighted regression model, MIC and Spearman's rank correlation coefficient to explore the predictor importance during feature selection step. These displays show the importance of comparing multiple association measures at once for different variable pairs.

Bibliography

A. Agresti. *Analysis of ordinal categorical data*, volume 656. John Wiley & Sons, 2010. [p4]




Figure 10: Conditional Association plot using linear layout. The display has variable pairs on the Y-axis and the value of association measures on the X-axis. The points corresponding to every variable pair represents the value of association measure for different levels of the conditioning variable and the overall value of association measure.

- A. Buja, A. M. Krieger, and E. I. George. A visualization tool for mining large correlation tables: The association navigator, 2016. [p2]
- D. Dua and C. Graff. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>. [p9]
- M. Friendly. Corrgrams: Exploratory displays for correlation matrices. *The American Statistician*, 56(4): 316–324, 2002. [p1, 2]
- S. Gerber. *scorr: s-CorrPlot: Visualizing Correlation*, 2022. URL <http://mckennapsean.com/scorrplot/>. R package version 1.0. [p2]
- M. Hills. On looking at large correlation matrices. *Biometrika*, 56(2):249–253, 1969. [p2]
- M. G. Kendall. The treatment of ties in ranking problems. *Biometrika*, 33(3):239–251, 1945. [p4]
- M. Kuhn, K. Johnson, et al. *Applied predictive modeling*, volume 26. Springer, 2013. [p16]
- M. Kuhn, S. Jackson, and J. Cimentada. *corrr: Correlations in R*, 2020. URL <https://CRAN.R-project.org/package=corrr>. R package version 0.4.3. [p2]
- P. Morgen and P. Biecek. *corrgrapher: Explore Correlations Between Variables in a Machine Learning Model*, 2020. URL <https://CRAN.R-project.org/package=corrgrapher>. R package version 1.0.4. [p2]
- D. J. Murdoch and E. Chow. A graphical display of large correlation matrices. *The American Statistician*, 50(2):178–180, 1996. [p2]
- U. Olsson. Maximum likelihood estimation of the polychoric correlation coefficient. *Psychometrika*, 44 (4):443–460, 1979. [p4]
- D. N. Reshef, Y. A. Reshef, H. K. Finucane, S. R. Grossman, G. McVean, P. J. Turnbaugh, E. S. Lander, M. Mitzenmacher, and P. C. Sabeti. Detecting novel associations in large data sets. *science*, 334(6062): 1518–1524, 2011. [p1, 3, 4]
- A. Samba. *linkspotter: Bivariate Correlations Calculation and Visualization*, 2020. URL <https://CRAN.R-project.org/package=linkspotter>. R package version 1.3.0. [p2]
- G. J. Székely, M. L. Rizzo, and N. K. Bakirov. Measuring and testing dependence by correlation of distances. *The annals of statistics*, 35(6):2769–2794, 2007. [p1, 3, 4]
- H. Theil. On the estimation of relationships involving qualitative variables. *American Journal of Sociology*, 76(1):103–154, 1970. [p4]
- E. R. Tufte. *The Visual Display of Quantitative Information*. Graphics Press, USA, 1986. ISBN 096139210X. [p1]
- J. W. Tukey and P. A. Tukey. Computer graphics and exploratory data analysis: An introduction. In *Proceedings of the sixth annual conference and exposition: computer graphics*, volume 85, pages 773–785, 1985. [p1, 16]
- T. Wei and V. Simko. *R package 'corrplot': Visualization of a Correlation Matrix*, 2021. URL <https://github.com/taiyun/corrplot>. (Version 0.92). [p2]
- H. Wickham, M. Averick, J. Bryan, W. Chang, L. D. McGowan, R. François, G. Golemund, A. Hayes, L. Henry, J. Hester, M. Kuhn, T. L. Pedersen, E. Miller, S. M. Bache, K. Müller, J. Ooms, D. Robinson, D. P. Seidel, V. Spinu, K. Takahashi, D. Vaughan, C. Wilke, K. Woo, and H. Yutani. Welcome to the tidyverse. *Journal of Open Source Software*, 4(43):1686, 2019. doi: 10.21105/joss.01686. [p2, 7]
- L. Wilkinson, A. Anand, and R. Grossman. Graph-theoretic scagnostics. In *Information Visualization, IEEE Symposium on*, pages 21–21. IEEE Computer Society, 2005. [p16]

Amit Chinwan
Maynooth University
Hamilton Institute
Maynooth, Ireland
amit.chinwan.2019@mumail.ie

Catherine Hurley
Maynooth University
Department of Mathematics and Statistics
Maynooth, Ireland
catherine.hurley@mu.ie